



Persona-Driven Document Intelligence

This project is part of **Challenge Round 1B** and focuses on analyzing multiple UPI fraud detection research documents using AI. The goal is to **automatically identify the most relevant sections** of provided research PDFs that help answer a specific job/task — such as determining effective machine learning algorithms for UPI fraud detection.

Project Structure

```

.
├── app/
│   ├── input/          # Input PDFs to be analyzed
│   ├── output/         # Output folder containing JSON results
│   ├── utils/          # Utility modules for PDF parsing, embedding, ranking, etc.
│   │   ├── embedder.py  # SentenceTransformer wrapper
│   │   ├── formatter.py # JSON result builder
│   │   ├── pdf_parser.py # Extracts text from PDF pages
│   │   └── ranker.py    # Ranks sections based on relevance
│   └── main.py          # Main script to run the analysis
├── test_input/
│   └── challenge1b_input.json # Contains metadata, persona, job, and document list
├── output/
│   └── challenge1b_output.json # Final generated output
├── approach_explanation.md # [Optional] Describes approach taken (not executed)
├── Dockerfile            # Docker setup for containerized execution
└── requirements.txt      # Python dependencies

```

Problem Statement

> **Task**: Build an intelligent document analysis system that extracts and ranks the top 5 most relevant sections from a set of academic PDFs, customized for a specific persona and their job-to-be-done.

Use Case

- **Persona**: Fraud Analyst at a FinTech company
- **Job-To-Be-Done**: Identify the most effective **machine learning algorithms and techniques** for detecting **UPI (Unified Payments Interface) fraud**.

Key Requirements

- The system should analyze a **collection of academic PDFs**.
- It must extract content at the **section level** (not just paragraphs or keywords).
- The top **5 most relevant sections** should be selected and ranked based on their alignment with the persona’s job.
- The analysis must be **persona-aware**, meaning it should account for the background, goals, and needs of the specific role (Fraud Analyst).
- The relevance should be evaluated based on **semantic context**, not just keyword matching.

How It Works

The pipeline performs the following steps:

1. Input Loading

- Reads metadata from `challenge1b_input.json`, which includes:
 - List of documents (PDFs)
 - Persona (user role)
 - Job/task description

2. PDF Parsing (`utils/pdf_parser.py`)

- Uses PyMuPDF to extract text from each page of each PDF.
- Stores each page as a section with:
 - Page number
 - Text content
 - Title: `Page X`

3. Embedding (`utils/embedder.py`)

- Uses `sentence-transformers` model (`all-MiniLM-L6-v2`) to convert texts into embeddings.
- Also embeds the job/task description.

4. Ranking (`utils/ranker.py`)

- Computes cosine similarity between the query embedding and section embeddings.
- Selects top 5 sections with the highest similarity.

5. Formatting Output (`utils/formatter.py`)

- Builds a JSON output with:
- Metadata
- Ranked extracted sections (titles, page numbers)
- Subsection analysis with top 2000 characters from each section

6. Output Saving

- Result is saved to `app/output/challenge1b_output.json`.

Dockerized Setup

Run the pipeline using Docker:

Build the Image:

```
```bash
docker build -t upi-fraud-analyzer .
```
```

Run the Container:

```
```bash
docker run --rm -v ${pwd}/app/output:/app/app/output upi-fraud-analyzer
```
```

This mounts the output folder so results persist outside the container.

Requirements

Install dependencies locally (if not using Docker):

```
```bash
pip install -r requirements.txt
```
```

Key dependencies:

- `sentence-transformers`
- `torch`
- `PyMuPDF`

Input Example (`test_input/challenge1b_input.json`)

```
```json
{
 "persona": {
 "role": "Fraud Analyst at a FinTech company"
 },
 "job_to_be_done": {
 "task": "Identify the most effective machine learning algorithms, preprocessing methods, and evaluation techniques to build a real-time UPI fraud detection system."
 },
 "documents": [
 {
 "filename": "Comparative_Analysis_of_UPI_Fraud_Detection_Using_Ensemble_Learning.pdf",
 "title": "Comparative Analysis of UPI Fraud Detection Using Ensemble Learning"
 },
 {
 "filename": "Fraud_Fighters_-_How_AI_and_ML_are_Revolutionizing_UPI_Security.pdf",
 "title": "Fraud Fighters - How AI and ML are Revolutionizing UPI Security"
 },
 {
 "filename": "Secure_UPI_Machine_Learning-Driven_Fraud_Detection_System_for_UPI_Transactions.pdf",
 "title": "Secure UPI - ML Driven Fraud Detection System"
 }
]
}
```
```

Output Format (output/challenge1b_output.json)

```
```json
{
 "metadata": {
 "input_documents": [...],
 "persona": "...",
 "job_to_be_done": "...",
 "processing_timestamp": "YYYY-MM-DDTHH:MM:SS"
 },
 "extracted_sections": [
 {
 "document": "...",
 "section_title": "...",
 "importance_rank": 1,
 "page_number": 3
 }
],
 "subsection_analysis": [
 {
 "document": "...",
 "refined_text": "Top 2000 chars...",
 "page_number": 3
 }
]
}
```

---

## ✔ Summary

- Multi-document processor for UPI fraud literature
- Uses semantic embeddings to rank sections by relevance
- Persona-aware, job-guided information retrieval
- Fully Dockerized, portable, and reproducible

---